

Evaluation of NoSQL databases for large-scale decentralized microblogging



Cassandra & Couchbase



Alexandre Fonseca, Anh Thu Vu, Peter Grman

Decentralized Systems - 2nd semester 2012/2013
Universitat Politècnica de Catalunya

Microblogging

- Blogging with very small & concise posts.
- Users subscribe to others by following them.
- Real-time interactions between growing userbase.
- Useful for:
 - Real-time updates & news.
 - Marketing and public relations campaigns.
 - Research into social interaction & perception.
- Examples:
 - Twitter, Tumblr, identi.ca,
 - Facebook & Google+ status updates.



Microblogging - decentralization

- Twitter:

- Over 500 million registered users.
- 340 million tweets per day.
- More than 12TB of data per day in 2010.



- Hard to provide this service relying solely on a small number of centralized servers.
 - Scaling up has its limits.

Microblogging - decentralization

- 2 possible ways to decentralize:
 - Decentralize dedicated servers managed by the service provider - better load sharing.
 - Allow users to share the system load between them (similar to voluntary computing).



Relational Database Systems (RDS)

- Typical enterprise and service data storage.
- Tables and table rows as storage units.
- Highly normalized.
- Storage divided into entity types.
 - One table per entity type.
 - One row per entity instance.
 - Many-to-many relationship tables.
- Transaction support for maximum consistency.


RDS - Twitter Example

Users		
ID	Name	Password
1	Alex	*****
2	Casey	*****
3	Peter	*****

Tweets			
ID	PosterID	Time	Body
1	1	1345453	Burguer
2	2	1345455	Nutella
3	3	1345457	Beer

Followers	
ID	Follows
1	2
1	3
2	1

Timeline for User Alex			
Tweet ID	Poster	Time	Body
2	Casey	1345453	Nutella
3	Peter	1345455	Beer



NoSQL

- New (or is it?) data storage paradigm:
 - KISS.
 - Simple key-value stores in its essence.
 - (Usually) no transaction control or strict schemas.
 - Incentive denormalization.
 - Let the application worry about business logic and entity relationships.
 - Focus on:
 - Flexibility (of schemas, topologies, configurations)
 - Scalability
 - Robustness
 - Performance

Not
Only SQL

NoSQL - Twitter Example

Users

```
[{'id': 1, 'name': 'Alex',  
'password': '***', 'following':[2, 3]},  
{ 'id': 2, 'name': 'Casey',  
'password': '***', 'following':[1]},  
{ 'id': 3, 'name': 'Peter',  
'password': '***', 'following':[1, 2]}
```

Tweets

```
[{'id': 1, 'posterID': 1,  
'time': 1345453, body: 'Burger'},  
{ 'id': 2, 'posterID': 2,  
'time': 1345455, body: 'Nutella'},  
{ 'id': 3, 'posterID': 3,  
'time': 1345457, body: 'Beer'}
```

Timeline-View

```
[{'user': 1, 'tweetID': 2, 'postedBy': 2, tweetBody: 'Nutella'},  
{ 'user': 1, 'tweetID': 3, 'postedBy': 3, tweetBody: 'Beer'},  
{ 'user': 2, 'tweetID': 1, 'postedBy': 1, tweetBody: 'Burger'},  
{ 'user': 3, 'tweetID': 1, 'postedBy': 1, tweetBody: 'Burger'},  
{ 'user': 3, 'tweetID': 2, 'postedBy': 2, tweetBody: 'Nutella'}
```


NoSQL - Twitter Example

Users

```
[{'id': 1, 'name': 'Alex',  
'password': '***', 'followers':[2, 3]},  
{ 'id': 2, 'name': 'Casey',  
'password': '***', 'followers':[1, 3]},  
{ 'id': 3, 'name': 'Peter',  
'password': '***', 'followers':[1]}]
```

Tweets

```
[{'id': 1, 'posterID': 1,  
'time': 1345453, body: 'Burger'},  
{ 'id': 2, 'posterID': 2,  
'time': 1345455, body: 'Nutella'},  
{ 'id': 3, 'posterID': 3,  
'time': 1345457, body: 'Beer'}]
```

Timeline-View

```
[{'user': 1, 'tweetID': 2, 'postedBy': 2, tweetBody: 'Nutella'},  
{ 'user': 1, 'tweetID': 3, 'postedBy': 3, tweetBody: 'Beer'},  
{ 'user': 2, 'tweetID': 1, 'postedBy': 1, tweetBody: 'Burger'},  
{ 'user': 3, 'tweetID': 1, 'postedBy': 1, tweetBody: 'Burger'},  
{ 'user': 3, 'tweetID': 2, 'postedBy': 2, tweetBody: 'Nutella'}]
```

Cassandra

- Initially developed by Facebook.
- Released in 2008.
- Apache Top-Project by 2010.
- Latest stable release (1.2.4, April 2013)
- Data model:
 - Hybrid between key-value and tabular storage.
 - Data split in column families (like RDS tables).
 - Column families split into rows (indexed by row keys).
 - Dynamic columns (schema).
- Interface: CQL via Thrift.



Cassandra

Cassandra

- All database nodes play the same role.
- Nodes divided into variable number of virtual nodes:
 - Organized into ring structure.
 - Assigned tokens determine data assignment:
 - Consistent hashing of partition key (1st column of primary key).
 - Hashing can be parametrized.
- Cluster can be given topology awareness:
 - Try to keep request resolution local to rack or datacenter.
 - Replicate data over different data centers and racks.
- Replication and consistency options all configurable.

Couchbase



Current release v2.0.1 (Apr 9, 2013)

Couchbase

- Key - Value
 - Key: string, no space allowed, 250 bytes limit
 - Value: JSON document, integer or serializable byte stream, 20MB limit
- View:
 - MapReduce: to index all data
 - Incrementally and periodically updated

Couchbase

- All nodes play the same role
- A bucket (or database) is
 - divided into small subsets, called vBuckets
 - distributed in the cluster.
- Replication factor:
 - configurable for each bucket with a max of 3.
- Consistency:
 - Strong consistency for access with "key"
 - Eventual consistency for queries on views

PyDLoader

- Custom Python script
- 2 components:
 - Manager:
 - Interactive console.
 - Deploy new nodes.
 - Install and control slaves on those nodes.
 - Slaves:
 - Automatic setup, populating and takedown of database clusters (database slaves).
 - Generation of application workload towards database clusters (workload slaves).
- Used libraries: Boto (AWS), Paramiko (SSH), RpyC (RPC).



Evaluation

- Done using Amazon Web Services (AWS).
- 6 database nodes (m1.small)
 - 1.7 GB of RAM
 - 1 compute unit
 - 150GB of storage
- 12 workload nodes (t1.micro)
 - 615 MB of RAM
 - 1 compute unit
 - No local storage (NAS)
- Distributed equally over 2 datacenters in Ireland.

Evaluation

- Focus on 3 main operations:
 - Tweet
 - Userline (all tweets by user)
 - Timeline (all tweets by people user is following).
- Limit to 50 items per userline/timeline.
- Basic data structure:
 - Aggregate data according to operations, not entities:
 - Timeline table/bucket:
 - UserID, TweetID, PostedBy, Body
 - Userline table/bucket:
 - UserID, TweetID, Body


Ease of Setup

- **Cassandra - Awesome!**
 - Install, configure addresses, partitioner, snitch, replication factors and seed nodes, launch!
 - Automatic partitioning and replication.
 - Automatic adjustment to churn.
- **Couchbase - Equivalently Awesome!**
 - Install, configure RAM, directory, launch!
 - Easily add/failover/remove servers
 - Rebalance: background process, asynchronous, incremental
 - Automatic replication, partitioning and node failure detection.

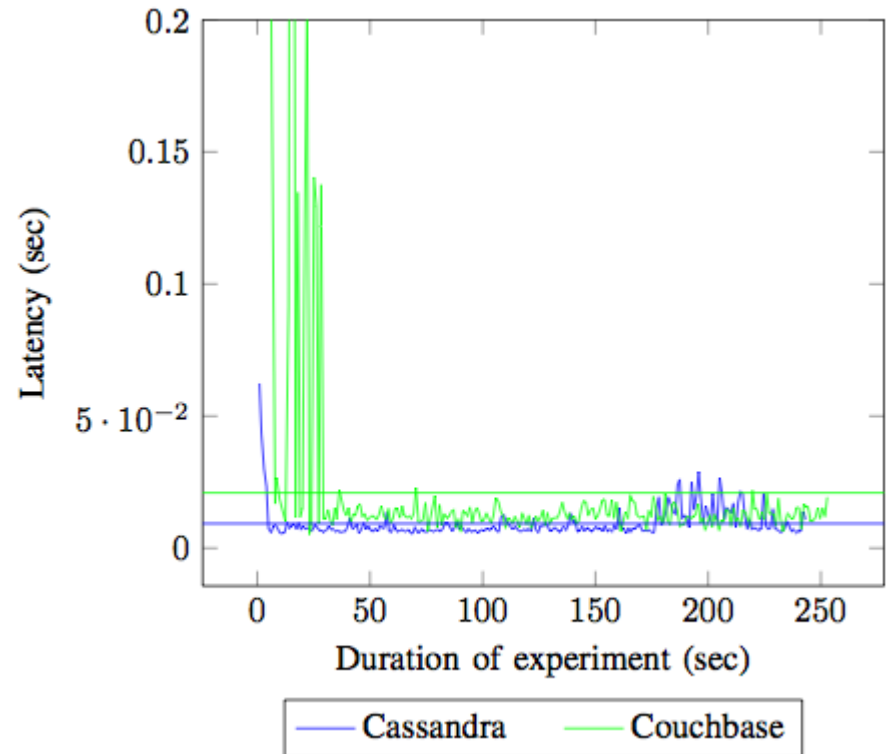
Setup time

- **Cassandra:**
 - 6 nodes form and stabilize ring after ~2 minutes.
 - Populating with sample data:
 - ~1 minute in 1 node configuration.
 - 6.5 minutes in 6 node configuration.
- **Couchbase:**
 - 6 nodes form and stabilize ring after ~2 minutes.
 - Populating with sample data:
 - ~1 minute in 1 node configuration.
 - ~2 minutes in 6 node configuration.

Latency

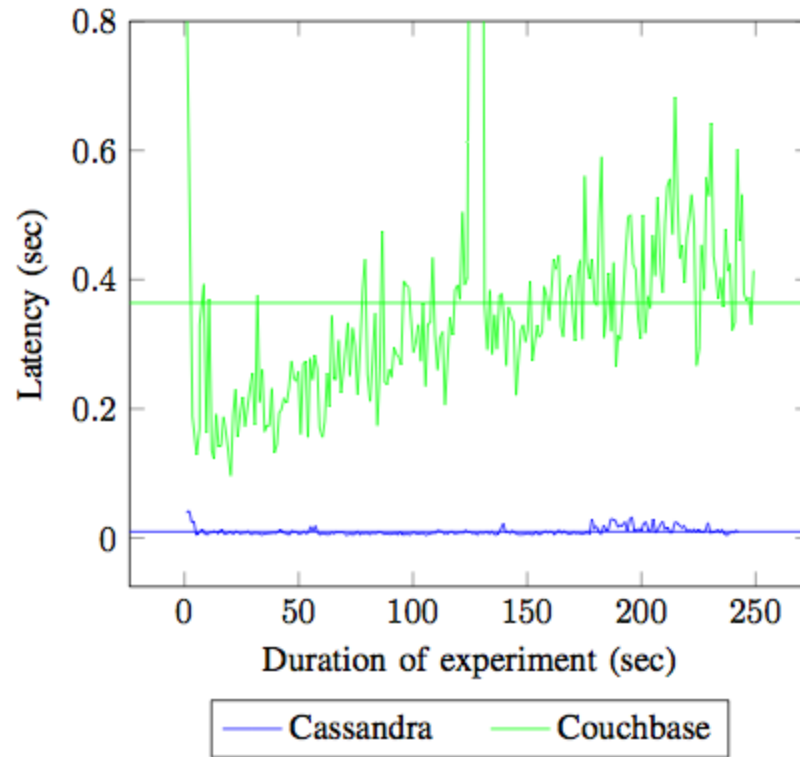


Tweet Latency




Userline Latency

Latency




Timeline Latency

Normalized latency

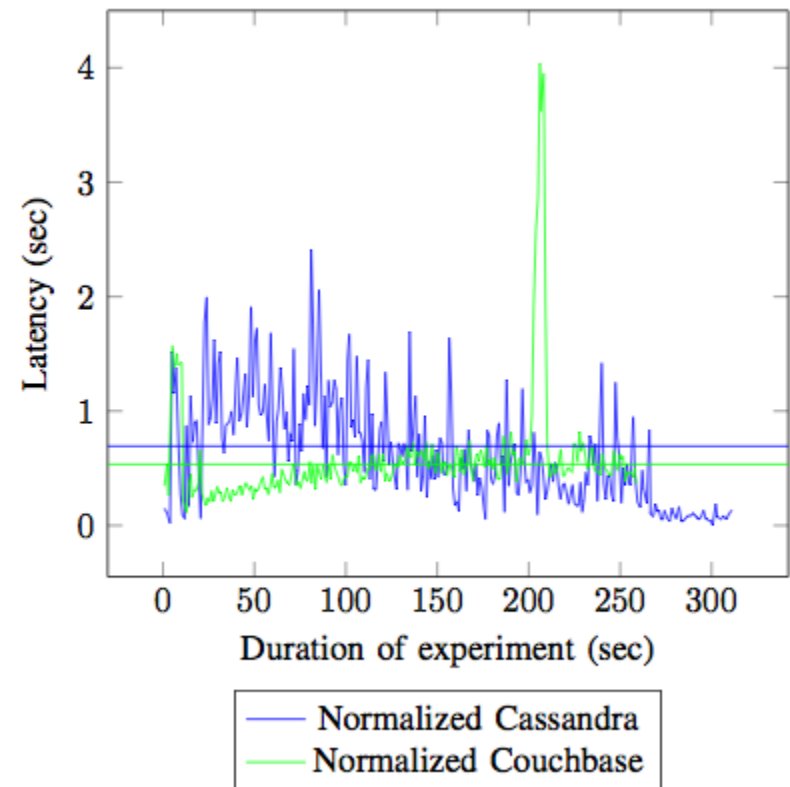


Tweet Latency

Normalized latency

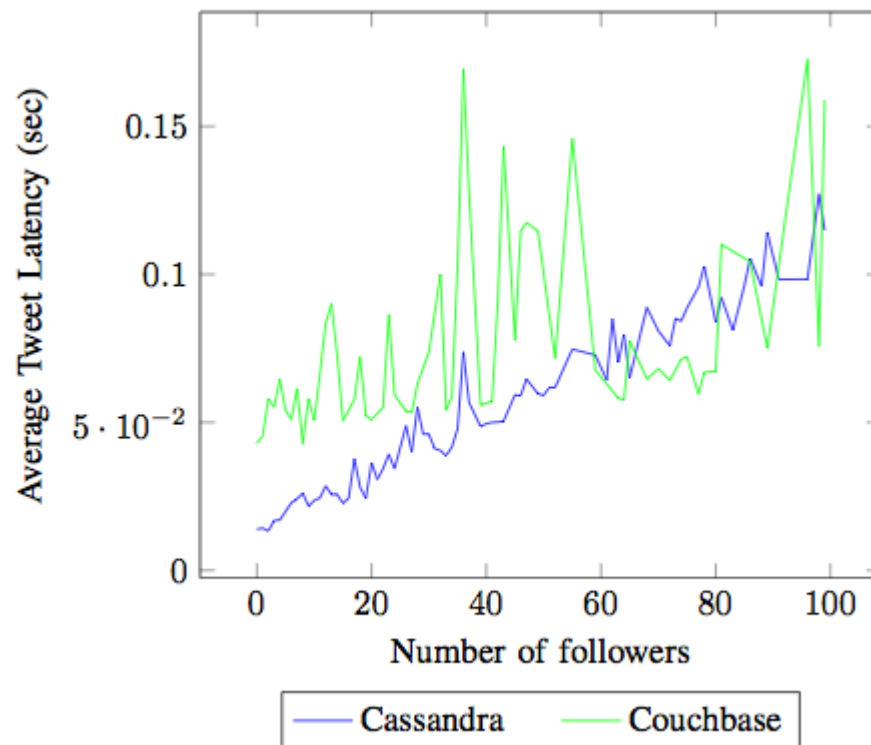


Userline Latency



Timeline Latency


Tweets & Denormalization



Tweet Latency v.s. Number of Followers

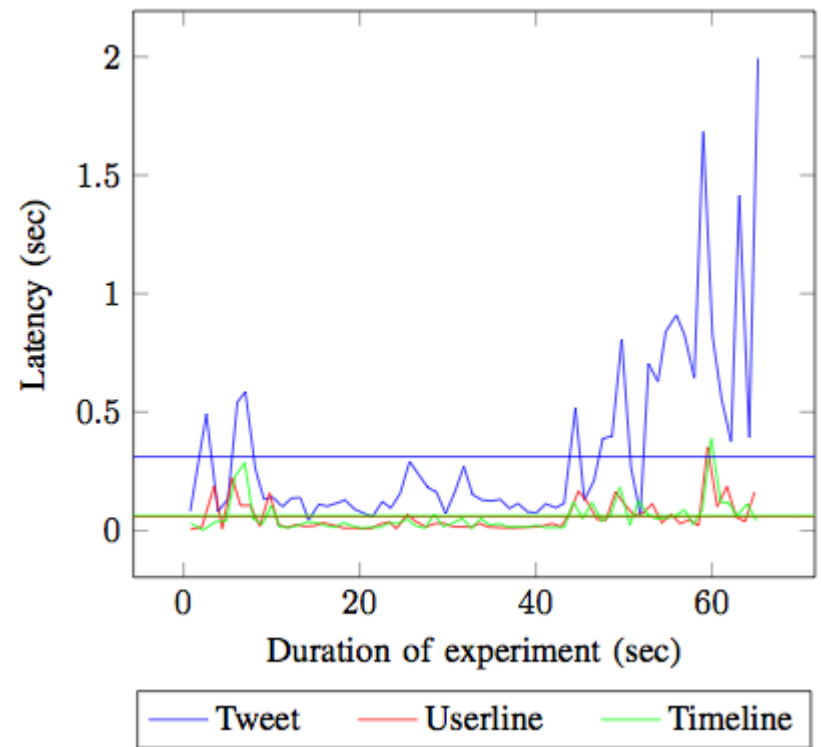
Tweets & Denormalization

- Immediate denormalization not scalable.
- How to make asynchronous?
 - Cassandra:
No native support.
External processing.
 - Couchbase: Views!



Reconfiguration Latency


- Cassandra
Adding a new node
to a cluster: ~5 mins.



Latency while node joining cluster

Reconfiguration Latency

- Couchbase
 - Adding a new node to a cluster: immediate
 - BUT rebalance ~ 30 mins




Latency while node joining cluster

Consistency/Convergence

- Cassandra:
 - Average of 0.096498 seconds to detect new tweet.
 - Standard deviation of 0.096319:
 - Same datacenter => very fast detection.
 - Different datacenter => slower detection.
- Couchbase:
 - Average of 0.001593 seconds to detect new tweet with standard deviation of 0.000526
 - The delay for new tweet to appear on timeline proportional to the schedule period

Replication


- Cassandra:
 - Very flexible in terms of replication configuration.
 - Per datacenter replication factors with no hard-coded limitations.



Behaviour under crash of 2 nodes
@ second 100th

Replication

- Couchbase
Automatic and configurable per bucket with a limit of (1+3) replicas.

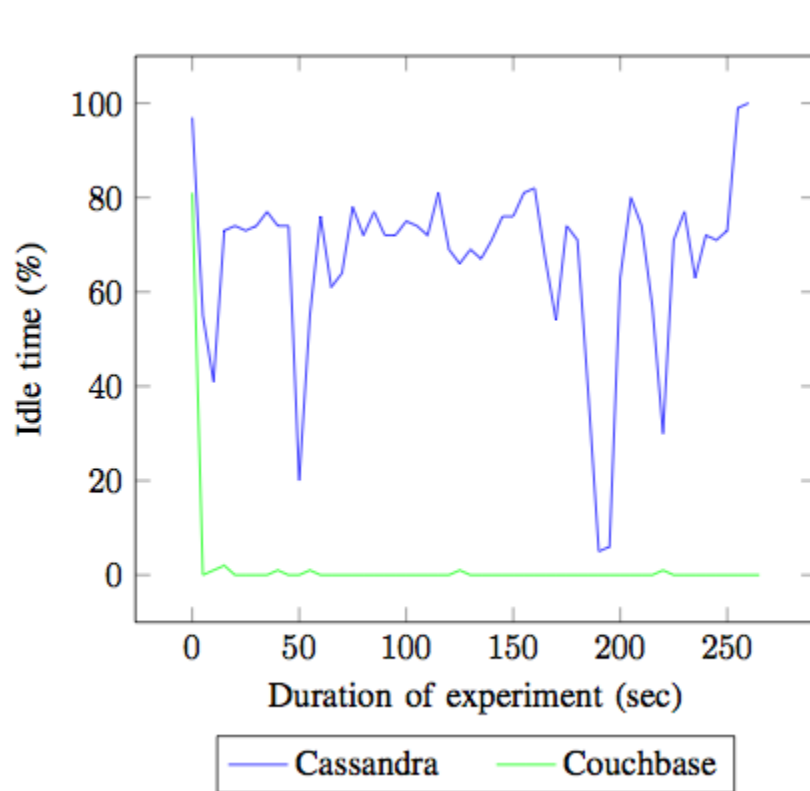


Behaviour under crash of 2 nodes
@ second 30th, 100th


Load Balancing

- **Cassandra:**
 - Average data ownership per node: 16.68%
 - Standard deviation: 1.23%
- **Couchbase:**
 - Evenly distributed with standard deviation:
 - 0.65% for data on disk
 - 0.04% for RAM usage

System Load

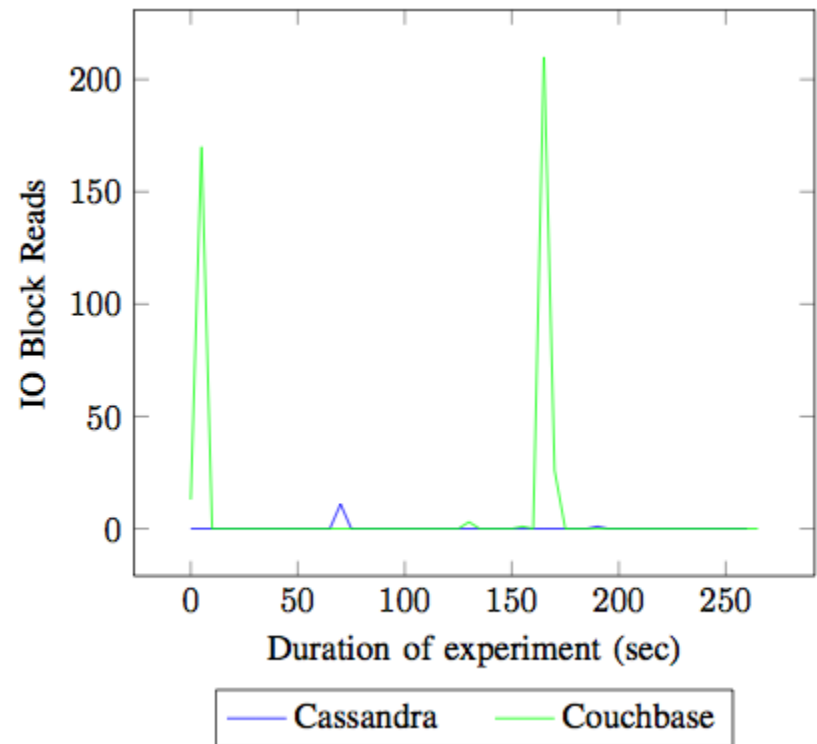



CPU Idle Time



Free Memory

System Load



Disk IO

Disk Space Usage

- SQLite database: 1.6MB
- Cassandra:
 - Fully denormalized (single node): 16.31MB.
 - Fully denormalized (6-nodes): 9.5MB/node.
 - Includes partitioning and replication.
 - "Normalized" (no body in userline/timeline):
 - Single node: 2.8MB
 - 6-nodes: 1.6MB/node.
 - Commit log after populating: 54MB
 - Need 50% of free disk space at all times:
 - Column family compactions.
 - Data redistributions.

Disk Space Usage

- Couchbase:
 - Total of 250MB distributed across 6 nodes.
 - No minimum free space requirement.
- Lack of disk usage limitations in both DBs:
 - Not ideal for voluntary computing systems.

Conclusion

- Easy cluster setup
 - Allows horizontal scaling over multiple nodes
- Improved performance through denormalization
- Higher storage requirements
- The future is hybrid
 - A mix of RDS and NoSQL-Systems
 - Cassandra's CQL, CouchBase's Views, NoSQL in RDS